

# Self-Driving Car Engineer Nanodegree

## Deep Learning

### Project: Build a Traffic Sign Recognition Classifier ¶

In this notebook, a template is provided for you to implement your functionality in stages which is required to successfully complete this project. If additional code is required that cannot be included in the notebook, be sure that the Python code is successfully imported and included in your submission, if necessary. Sections that begin with **'Implementation'** in the header indicate where you should begin your implementation for your project. Note that some sections of implementation are optional, and will be marked with **'Optional'** in the header.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a **'Question'** header. Carefully read each question and provide thorough answers in the following text boxes that begin with **'Answer:'**. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. In addition, Markdown cells can be edited by typically double-clicking the cell to enter edit mode.

## Imports

```
In [1]: import pickle
import time
import os
import scipy.ndimage

import numpy as np
import pandas as pd
import tensorflow as tf
from matplotlib.colors import ListedColormap
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import matplotlib.gridspec as gridspec

from sklearn import datasets
from sklearn.utils import shuffle
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score, f1_score
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder
from sklearn import decomposition
import cv2
```

## Supporting Functions

```
In [2]: #one hot coding function
def OHE_Encode(Y_tr,N_classes):
    OHC = OneHotEncoder()
    Y_ohc = OHC.fit(np.arange(N_classes).reshape(-1, 1))
    Y_labels = Y_ohc.transform(Y_tr.reshape(-1, 1)).toarray()
    return Y_labels

def OHE_Validate(cls,y):
    check = np.linalg.norm(np.argmax(cls,axis=1)-y)
    if check == 0:
        print('One hot encoding Validated')
    else:
        print('One hot encoding doesnt match the output, check code!!!')
```

---

## Step 0: Load The Data

```
In [3]: print('Tensor Flow version : '+tf.__version__)
training_file = 'traffic-signs-data/train.p'
testing_file = 'traffic-signs-data/test.p'

with open(training_file, mode='rb') as f:
    train = pickle.load(f)
with open(testing_file, mode='rb') as f:
    test = pickle.load(f)

X_train, y_train = train['features'], train['labels']
X_test, y_test = test['features'], test['labels']

assert(len(X_train)== len(y_train))
assert(len(X_test)== len(y_test))

#data_pd = pd.read_csv('signnames.csv')
#data_pd.head()
```

Tensor Flow version : 0.12.1

---

## Step 1: Dataset Summary & Exploration

The pickled data is a dictionary with 4 key/value pairs:

- 'features' is a 4D array containing raw pixel data of the traffic sign images, (num examples, width, height, channels).
- 'labels' is a 2D array containing the label/class id of the traffic sign. The file `signnames.csv` contains id -> name mappings for each id.
- 'sizes' is a list containing tuples, (width, height) representing the the original width and height the image.
- 'coords' is a list containing tuples, (x1, y1, x2, y2) representing coordinates of a bounding box around the sign in the image. **THESE COORDINATES ASSUME THE ORIGINAL IMAGE. THE PICKLED DATA CONTAINS RESIZED VERSIONS (32 by 32) OF THESE IMAGES**

Complete the basic data summary below.

```

In [4]: ### Replace each question mark with the appropriate value.

# TODO: Number of training examples
n_train = len(X_train)

# TODO: Number of testing examples.
n_test = len(X_test)

# TODO: What's the shape of an traffic sign image?
image_shape = X_train[0].shape

# TODO: How many unique classes/labels there are in the dataset.
n_classes = len(set(y_train))

# Number of Channels
n_channels = X_train[0].shape[2]

print("Number of training examples =", n_train)
print("Number of testing examples =", n_test)
print("Image data shape =", image_shape)
print('Training labels shape', y_train.shape)
print("Number of classes =", n_classes)
print("Number of channels =", n_channels)
print("Number of features =", X_train.shape[1])

#One Hot Encoding
y_train_OneHot = OHE_Encode(y_train,n_classes)
y_test_OneHot = OHE_Encode(y_test,n_classes)

OHE_Validate(y_test_OneHot ,y_test)
OHE_Validate(y_train_OneHot,y_train)

Number of training examples = 39209
Number of testing examples = 12630
Image data shape = (32, 32, 3)
Training labels shape (39209,)
Number of classes = 43
Number of channels = 3
Number of features = 32
One hot encoding Validated
One hot encoding Validated

```

Visualize the German Traffic Signs Dataset using the pickled file(s). This is open ended, suggestions include: plotting traffic sign images, plotting the count of each sign, etc.

The [Matplotlib](http://matplotlib.org/) (<http://matplotlib.org/>) [examples](http://matplotlib.org/examples/index.html) (<http://matplotlib.org/examples/index.html>) and [gallery](http://matplotlib.org/gallery.html) (<http://matplotlib.org/gallery.html>) pages are a great resource for doing visualizations in Python.

**NOTE:** It's recommended you start with something simple first. If you wish to do more, come back to it after you've completed the rest of the sections.

```
In [5]: print('Training Data')
train_features = np.array(train['features'])
train_labels = np.array(train['labels'])

inputs_per_class = np.bincount(train_labels)
max_inputs = np.max(inputs_per_class)


mpl_fig = plt.figure(figsize=(12,8))
ax = mpl_fig.add_subplot(111)
ax.set_ylabel('Inputs')
ax.set_xlabel('Class')
ax.set_title('Number of inputs per class')
ax.bar(range(len(inputs_per_class)), inputs_per_class, 1/3, color='blue', label='Inputs per class')
plt.show()

print('Testing Data')
test_features = np.array(test['features'])
test_labels = np.array(test['labels'])


inputs_per_class = np.bincount(test_labels)
max_inputs = np.max(inputs_per_class)

mpl_fig = plt.figure(figsize=(12,8))
ax = mpl_fig.add_subplot(111)
ax.set_ylabel('Inputs')
ax.set_xlabel('Class')
ax.set_title('Number of inputs per class')
ax.bar(range(len(inputs_per_class)), inputs_per_class, 1/3, color='blue', label='Inputs per class')
plt.show()
```

## Training Data



## Testing Data



```
In [6]: data_i = [[i,sum(y_train == i)] for i in range(len(np.unique(y_train)))]
data_i_sorted = sorted(data_i, key=lambda x: x[1])
data_pd = pd.read_csv('signnames.csv')
data_pd['Occurance'] = pd.Series(np.asarray(data_i_sorted).T[1], index=np.asar
ray(data_i_sorted).T[0])
data_pd_sorted = data_pd.sort_values(['Occurance'],ascending=
[0]).reset_index()
data_pd_sorted = data_pd_sorted.drop('index', 1)
data_pd_sorted
```


Out[6]:

	ClassId	SignName	Occurance
0	2	Speed limit (50km/h)	2250
1	1	Speed limit (30km/h)	2220
2	13	Yield	2160
3	12	Priority road	2100
4	38	Keep right	2070
5	10	No passing for vehicles over 3.5 metric tons	2010
6	4	Speed limit (70km/h)	1980
7	5	Speed limit (80km/h)	1860
8	25	Road work	1500
9	9	No passing	1470
10	7	Speed limit (100km/h)	1440
11	3	Speed limit (60km/h)	1410
12	8	Speed limit (120km/h)	1410
13	11	Right-of-way at the next intersection	1320
14	35	Ahead only	1200
15	18	General caution	1200
16	17	No entry	1110
17	31	Wild animals crossing	780
18	14	Stop	780
19	33	Turn right ahead	689
20	15	No vehicles	630
21	26	Traffic signals	600
22	28	Children crossing	540
23	23	Slippery road	510
24	30	Beware of ice/snow	450
25	16	Vehicles over 3.5 metric tons prohibited	420
26	34	Turn left ahead	420
27	6	End of speed limit (80km/h)	420
28	36	Go straight or right	390
29	22	Bumpy road	390
30	40	Roundabout mandatory	360
31	20	Dangerous curve to the right	360



	ClassId	SignName	Occurance
32	21	Double curve	330
33	39	Keep left	300
34	29	Bicycles crossing	270
35	24	Road narrows on the right	270
36	41	End of no passing	240
37	42	End of no passing by vehicles over 3.5 metric ...	240
38	32	End of all speed and passing limits	240
39	27	Pedestrians	240
40	37	Go straight or left	210
41	19	Dangerous curve to the left	210
42	0	Speed limit (20km/h)	210

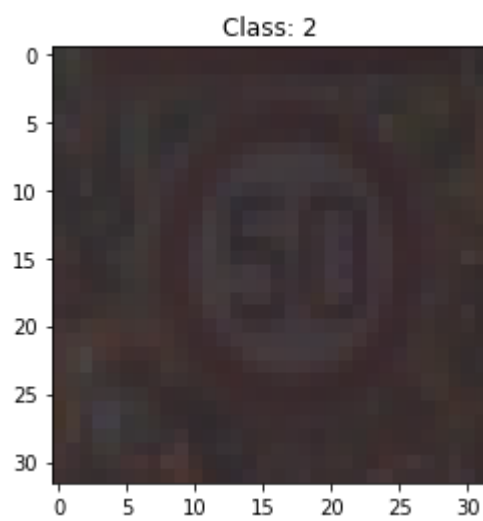

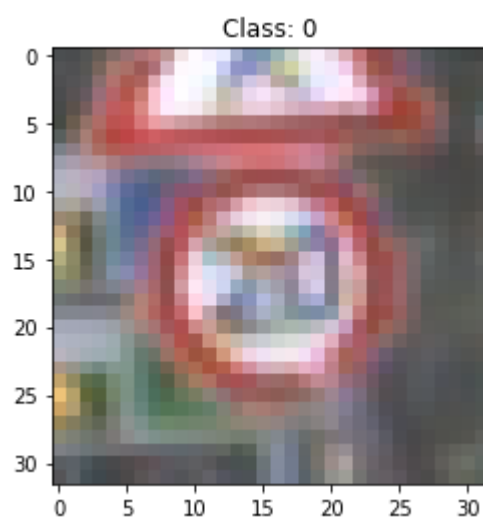
```
In [7]: plt.figure(figsize=(12,8))
plt.bar(range(n_classes),height=data_pd_sorted["Occurance"])
plt.show()
```




```
In [8]: sign_dict = {}

print('List of Classes')
mpl_fig = plt.figure()
for i in range(n_classes):
    for j in range(len(train_labels)):
        if (i == train_labels[j]):
            sign_dict[i]=train_features[j]
            plt.title('Class: ' + str(i))
            plt.imshow(train_features[j])
            plt.show()
            break
```

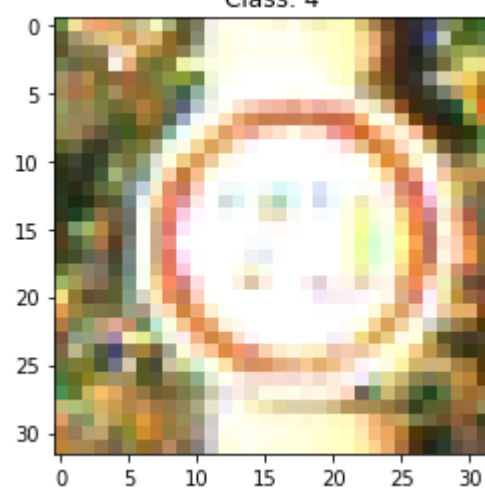
## List of Classes




Class: 3




Class: 4




Class: 5



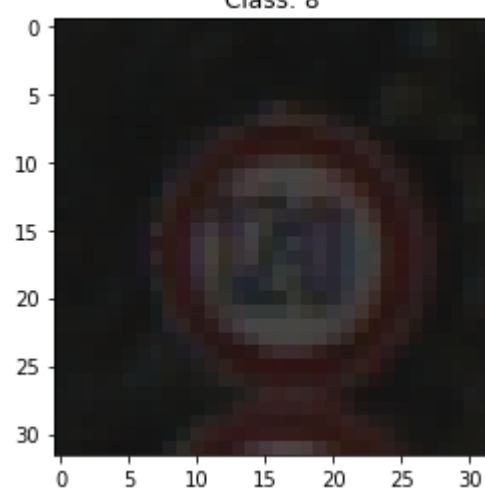
Class: 6




Class: 7




Class: 8




Class: 9




Class: 10




Class: 11




Class: 12




Class: 13



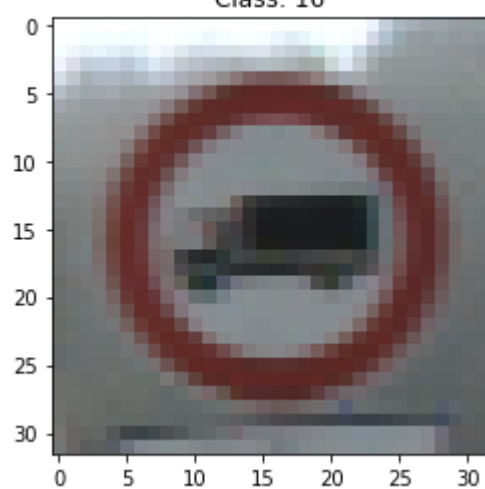
Class: 14




Class: 15



Class: 16




Class: 17







Class: 18




Class: 19




Class: 20




Class: 21




Class: 22




Class: 23




Class: 24




Class: 25




Class: 26




Class: 27






Class: 28




Class: 29







Class: 33




Class: 34




Class: 35




Class: 36




Class: 37




Class: 38




Class: 39




Class: 40



Class: 41







**Add Extra Data to avoid OverFitting**

```

In [9]: def get_index_dict(y_train):
        # Returns indices of each label
        # Assumes that the labels are 0 to N-1
        dict_indices = {}
        ind_all = np.arange(len(y_train))

        for i in range(len(np.unique(y_train))):
            ind_i = ind_all[y_train == i]
            dict_indices[i] = ind_i
        return dict_indices

def transform_image(image,ang_range,shear_range,trans_range):

    # Rotation

    ang_rot = np.random.uniform(ang_range)-ang_range/2
    rows,cols,ch = image.shape
    Rot_M = cv2.getRotationMatrix2D((cols/2,rows/2),ang_rot,1)

    # Translation
    tr_x = trans_range*np.random.uniform()-trans_range/2
    tr_y = trans_range*np.random.uniform()-trans_range/2
    Trans_M = np.float32([[1,0,tr_x],[0,1,tr_y]])

    # Shear
    pts1 = np.float32([[5,5],[20,5],[5,20]])

    pt1 = 5+shear_range*np.random.uniform()-shear_range/2
    pt2 = 20+shear_range*np.random.uniform()-shear_range/2

    pts2 = np.float32([[pt1,5],[pt2,pt1],[5,pt2]])

    shear_M = cv2.getAffineTransform(pts1,pts2)

    image = cv2.warpAffine(image,Rot_M,(cols,rows))

```

```

        image = cv2.warpAffine(image, Trans_M, (cols, rows))
        image = cv2.warpAffine(image, shear_M, (cols, rows))

    return image

def gen_extra_data(X_train, y_train, N_classes, n_each, ang_range, shear_range, trans_range, randomize_Var):
    dict_indices = get_index_dict(y_train)
    n_class = len(np.unique(y_train))
    X_arr = []
    Y_arr = []
    n_train = len(X_train)
    for i in range(n_train):
        for i_n in range(n_each):
            img_trf = transform_image(X_train[i], ang_range, shear_range, trans_range)
            X_arr.append(img_trf)
            Y_arr.append(y_train[i])

    return X_arr, Y_arr

i_train = 1
ang_rot = 10*0.9**(i_train)
trans_rot = 2*0.9**(i_train)
shear_rot = 2*0.9**(i_train)

#X_train, y_train = gen_extra_data(X_train, y_train, 43, 5, ang_rot, trans_rot, shear_rot, 1)

#print('Training Data')
#train_features = np.array(X_train)
#train_labels = np.array(y_train)

#inputs_per_class = np.bincount(train_labels)
#max_inputs = np.max(inputs_per_class)

#mpl_fig = plt.figure(figsize=(12, 8))
#ax = mpl_fig.add_subplot(111)
#ax.set_ylabel('Inputs')
#ax.set_xlabel('Class')
#ax.set_title('Number of inputs per class')
#ax.bar(range(len(inputs_per_class)), inputs_per_class, 1/3, color='blue', label='Inputs per class')
#plt.show()

```

---

## Step 2: Design and Test a Model Architecture

Design and implement a deep learning model that learns to recognize traffic signs. Train and test your model on the German Traffic Sign Dataset (<http://benchmark.ini.rub.de/?section=gtsrb&subsection=dataset>).

There are various aspects to consider when thinking about this problem:

- Neural network architecture
- Play around preprocessing techniques (normalization, rgb to grayscale, etc)
- Number of examples per label (some have more than others).
- Generate fake data.

Here is an example of a published baseline model on this problem

(<http://yann.lecun.com/exdb/publis/pdf/sermanet-ijcnn-11.pdf>). It's not required to be familiar with the approach used in the paper but, it's good practice to try to read papers like these.

**NOTE:** The LeNet-5 implementation shown in the classroom

(<https://classroom.udacity.com/nanodegrees/nd013/parts/fbf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/601ae704-1035-4287-8b11-e2c2716217ad/concepts/d4aca031-508f-4e0b-b493-e7b706120f81>) at the end of the CNN lesson is a solid starting point. You'll have to change the number of classes and possibly the preprocessing, but aside from that it's plug and play!

## Implementation

Use the code cell (or multiple code cells, if necessary) to implement the first step of your project. Once you have completed your implementation and are satisfied with the results, be sure to thoroughly answer the questions that follow.

## 2.1 Preprocessing

```

In [10]: ### Step 1
### Images are already resized the images to 32x32. so no padding required

### Step 2
### its important to shuffle data because order of data could effect how well
is network gets trained
X_train, y_train = shuffle(X_train, y_train, random_state=42)

### Step 3 Normalisation (not standardising)
# The goal is to independently normalize each feature component to the [0,1] range
# In image processing, normalization is a process that changes the range of pixel intensity values.
# In stochastic gradient descent, Normalization can sometimes improve the convergence speed of the algorithm

X_train_org = X_train
X_test_org = X_test
X_train = (X_train - X_train.mean()) / (np.max(X_train) - np.min(X_train))
X_test = (X_test - X_test.mean()) / (np.max(X_test) - np.min(X_test))
print('')
print('---Before and after Normalisation Sample---')
image_index = 2
plt.subplot(2,2,1)
plt.imshow(X_train_org[image_index])
plt.show()
plt.subplot(2,2,2)
plt.imshow(X_train[image_index])
plt.show()



### Step 4 Segmenting data into training, test, and validation
X_train, X_validation, y_train, y_validation = train_test_split(X_train, y_train, test_size=0.2, random_state=42)
y_validation_OneHot = OHE_Encode(y_validation,n_classes)
OHE_Validate(y_validation_OneHot,y_validation)

print('')
print('---Data shape---')
print('Validation data shape', X_validation.shape)
print('Training data shape', X_train.shape)
print('Testing data shape', X_test.shape)

print('')
print('---Label shape---')
print('Validation Label data shape', y_validation.shape)
print('Training Label data shape', y_train.shape)
print('Testing Label data shape', y_test.shape)

```

---Before and after Normalisation Sample---



One hot encoding Validated

---Data shape---

Validation data shape (7842, 32, 32, 3)

Training data shape (31367, 32, 32, 3)

Testing data shape (12630, 32, 32, 3)

---Label shape---

Validation Label data shape (7842,)

Training Label data shape (31367,)

Testing Label data shape (12630,)

## Step 3: Define Network

### 3.1 Network parameters

```
In [11]: n_filters = 32
kernel_size = (3, 3)
n_fc1 = 512
n_fc2 = 128
pool_size = 2 # i.e. (2,2)

dropout_conv = 0.9
dropout_fc = 0.9

weights_stddev = 0.1
weights_mean = 0.0
biases_mean = 0.0

padding = 'VALID'
if padding == 'SAME':
    conv_output_length = 6
elif padding == 'VALID':
    conv_output_length = 5
else:
    raiseException("Unknown padding.")
```

## 3.2 tf Graph input

```
In [12]: x_unflattened = tf.placeholder("float", [None, image_shape[0],image_shape[1],i
image_shape[2]])
x = x_unflattened

y_rawlabels = tf.placeholder("int32", [None])
y = tf.one_hot(y_rawlabels, depth=n_classes, on_value=1., off_value=0.,
axis=-1)
```

## 3.3 Model

```

In [13]: def weight_variable(shape, weight_mean, weight_stddev):
    return tf.Variable(tf.truncated_normal(shape, stddev=weight_stddev, mean=weight_mean))

def bias_variable(shape, bias_mean):
    return tf.Variable(tf.constant(bias_mean, shape=shape))

def conv2d(x, W, b, strides=3):
    x = tf.nn.conv2d(x, W, strides=[1, strides, strides, 1], padding='SAME')
    x = tf.nn.bias_add(x, b)
    return tf.nn.relu(x)

def maxpool2d(x, k=2, padding_setting='SAME'):
    return tf.nn.max_pool(x, ksize=[1, k, k, 1], strides=[1, k, k, 1],
padding=padding_setting)

def plot_metric_per_epoch(accuracies):
    x_epochs = []
    y_epochs = []
    for i, val in enumerate(accuracies):
        x_epochs.append(i)
        y_epochs.append(val)

    plt.figure(figsize=(15,8))
    plt.xlabel('epoch')
    plt.ylabel('score')
    plt.title('Score per epoch')
    #plt.legend()
    plt.grid()
    plt.scatter(x_epochs, y_epochs,s=50,c='lightgreen', marker='s', label='score')
    plt.show()

```

```

In [14]: weights = {
    'conv1': weight_variable([kernel_size[0], kernel_size[1], n_channels, n_filters], weights_mean, weights_stddev),
    'fc1': weight_variable([n_filters * conv_output_length**2, n_fc1], weights_mean, weights_stddev),
    'fc2': weight_variable([n_fc1, n_fc2], weights_mean, weights_stddev),
    'out': weight_variable([n_fc2, n_classes], weights_mean, weights_stddev)
}

biases = {
    'conv1': bias_variable([n_filters], biases_mean),
    'fc1': bias_variable([n_fc1], biases_mean),
    'fc2': bias_variable([n_fc2], biases_mean),
    'out': bias_variable([n_classes], biases_mean)
}

```



```
In [15]: def conv_net(model_x, model_weights, model_biases, model_pool_size,
                model_dropout_conv, model_dropout_fc, padding='SAME'):

    # Convolution Layer 1
    conv1 = conv2d(model_x, model_weights['conv1'], model_biases['conv1'])

    # Max Pooling (down-sampling)
    conv1 = maxpool2d(conv1, k=model_pool_size, padding_setting=padding)
    conv1 = tf.nn.dropout(conv1, model_dropout_conv)

    # Reshape conv1 output to fit fully connected layer input
    conv1_shape = conv1.get_shape().as_list()
    fc1 = tf.reshape(conv1, [-1,
conv1_shape[1]*conv1_shape[2]*conv1_shape[3]])

    # Fully connected Layer 1
    fc1 = tf.add(tf.matmul(fc1, model_weights['fc1']), model_biases['fc1'])
    fc1 = tf.nn.elu(fc1)
    fc1 = tf.nn.dropout(fc1, model_dropout_fc)

    # Fully connected Layer 2
    fc2 = tf.add(tf.matmul(fc1, model_weights['fc2']), model_biases['fc2'])
    fc2 = tf.nn.elu(fc2)
    fc2 = tf.nn.dropout(fc2, model_dropout_fc)

    # Output Layer
    output = tf.add(tf.matmul(fc2, model_weights['out']), model_biases['out'])

    return output
```

## Step 4: Train Model

### 4.1 Training Parameters

```
In [16]: #Learning Rate
#It represents the step that is taken in a gradient decent algorithm to find a
n optimal solution.
#If steps is too big, algorithm can go from peak to peak, and skip lows altoget
her.
#if steps is too small, algorithm will take lot of time to converge.
#To train a model, it is often recommended to lower the learning rate as the t
raining progresses.
#We will use exponential_decay function to do so.
#It requires 'initial learning rate' and 'global_step' value to compute the de
cayed learning rate.
learning_rate = 0.001
initial_learning_rate = learning_rate
annealing_rate = 1

#epochs number
#this represents the number of times that we will run our main "for" loop. thi
s is the loop that we will run to provide
#data to our algorithms for training and testing.
#this is not number of batches
training_epochs = 150

#Batch Size
#Deep learning algorithms are iterative in the sense that they load samples in
batches to avoid running out of memory.
#number of batch (bin) = number of training samples / batch size.
batch_size = 100

display_step = 1
anneal_mod_frequency = 15
print_accuracy_mod_frequency = 1
```

## 4.2 Model/Loss/Optimizer

```
In [17]: # Construct model
pred = conv_net(x, weights, biases, pool_size, dropout_conv, dropout_fc, paddi
ng=padding)
pred_probs = tf.nn.softmax(pred)

# Define Loss and optimizer
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(pred, y))
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
```

## 4.3 Train

```
In [18]: # Function to initialise the variables  
init = tf.global_variables_initializer()  
  
# Launch the graph  
sess = tf.Session()
```

```

# Initialise variables
sess.run(init)

# Initialise time logs
init_time = time.time()
epoch_time = init_time

five_epoch_moving_average = 0.
epoch_accuracies = []
print_accuracy_mod_frequency = 1
total_batch = int(len(X_train) / batch_size)

# Training cycle
for epoch in range(training_epochs):
    if five_epoch_moving_average > 0.96:
        break

    for i in range(total_batch):
        batch_x, batch_y = np.array(X_train[i * batch_size:(i + 1) * batch_size]), \
                               np.array(y_train[i * batch_size:(i + 1) * batch_size])
        _, epoch_cost = sess.run([optimizer, cost], feed_dict={x_unflattened:
batch_x, y_rawlabels: batch_y})

        # Display logs per epoch step
        if epoch % display_step == 0:
            last_epoch_time = epoch_time
            epoch_time = time.time()
            print("Epoch:", '%04d' % (epoch + 1), "cost:", "{:.9f}".format(epoch_cost), "Time since last epoch: ", epoch_time - last_epoch_time)

        # Anneal Learning rate
        if (epoch + 1) % anneal_mod_frequency == 0:
            learning_rate *= annealing_rate
            print("New learning rate: ", learning_rate)

        if (epoch + 1) % print_accuracy_mod_frequency == 0:
            correct_prediction = tf.equal(tf.argmax(pred, 1), tf.argmax(y, 1))
            accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
            # Line below needed only when not using `with tf.Session() as sess`
            with sess.as_default():
                epoch_accuracy = accuracy.eval({x_unflattened: X_validation, y_rawlabels: y_validation})
                epoch_accuracies.append(epoch_accuracy)

            ji = (len(epoch_accuracies), 5)[len(epoch_accuracies)>4]
            five_epoch_moving_average = np.sum(epoch_accuracies[epoch+1-ji:epoch+1])/ji
            print("epoch Accuracy (validation) = " + str(epoch_accuracy) + ",
five_epoch_moving_average = " + "{:.6f}".format(five_epoch_moving_average))

print("Optimization Finished!")
plot_metric_per_epoch(epoch_accuracies)

```

Epoch: 0001 cost: 0.835354030 Time since last epoch: 5.350874185562134  
epoch Accuracy (validation) = 0.719077, five\_epoch\_moving\_average = 0.719077  
Epoch: 0002 cost: 0.600972295 Time since last epoch: 5.97432804107666  
epoch Accuracy (validation) = 0.799413, five\_epoch\_moving\_average = 0.759245  
Epoch: 0003 cost: 0.408585429 Time since last epoch: 5.97832989692688  
epoch Accuracy (validation) = 0.838307, five\_epoch\_moving\_average = 0.785599  
Epoch: 0004 cost: 0.394816637 Time since last epoch: 5.910281181335449  
epoch Accuracy (validation) = 0.849911, five\_epoch\_moving\_average = 0.801677  
Epoch: 0005 cost: 0.222693145 Time since last epoch: 5.989338397979736  
epoch Accuracy (validation) = 0.865085, five\_epoch\_moving\_average = 0.814359  
Epoch: 0006 cost: 0.241149440 Time since last epoch: 5.948308706283569  
epoch Accuracy (validation) = 0.868783, five\_epoch\_moving\_average = 0.844300  
Epoch: 0007 cost: 0.158790365 Time since last epoch: 5.961318254470825  
epoch Accuracy (validation) = 0.885106, five\_epoch\_moving\_average = 0.861438  
Epoch: 0008 cost: 0.151555955 Time since last epoch: 5.985335350036621  
epoch Accuracy (validation) = 0.885871, five\_epoch\_moving\_average = 0.870951  
Epoch: 0009 cost: 0.166201204 Time since last epoch: 5.971325397491455  
epoch Accuracy (validation) = 0.900025, five\_epoch\_moving\_average = 0.880974  
Epoch: 0010 cost: 0.159873873 Time since last epoch: 5.992340803146362  
epoch Accuracy (validation) = 0.901556, five\_epoch\_moving\_average = 0.888268  
Epoch: 0011 cost: 0.203673854 Time since last epoch: 5.974327325820923  
epoch Accuracy (validation) = 0.910865, five\_epoch\_moving\_average = 0.896685  
Epoch: 0012 cost: 0.091472618 Time since last epoch: 5.9483091831207275  
epoch Accuracy (validation) = 0.913797, five\_epoch\_moving\_average = 0.902423  
Epoch: 0013 cost: 0.156288356 Time since last epoch: 5.9663214683532715  
epoch Accuracy (validation) = 0.921194, five\_epoch\_moving\_average = 0.909487  
Epoch: 0014 cost: 0.122915223 Time since last epoch: 5.96732234954834  
epoch Accuracy (validation) = 0.922724, five\_epoch\_moving\_average = 0.914027  
Epoch: 0015 cost: 0.100420021 Time since last epoch: 6.0033485889434814  
New learning rate: 0.001  
epoch Accuracy (validation) = 0.91673, five\_epoch\_moving\_average = 0.917062  
Epoch: 0016 cost: 0.110888556 Time since last epoch: 6.00434947013855  
epoch Accuracy (validation) = 0.924637, five\_epoch\_moving\_average = 0.919816  
Epoch: 0017 cost: 0.066086076 Time since last epoch: 6.013355731964111  
epoch Accuracy (validation) = 0.927442, five\_epoch\_moving\_average = 0.922545  
Epoch: 0018 cost: 0.085709855 Time since last epoch: 5.984334707260132  
epoch Accuracy (validation) = 0.927825, five\_epoch\_moving\_average = 0.923871  
Epoch: 0019 cost: 0.131604791 Time since last epoch: 6.001347064971924  
epoch Accuracy (validation) = 0.930885, five\_epoch\_moving\_average = 0.925504  
Epoch: 0020 cost: 0.109586865 Time since last epoch: 6.029367208480835  
epoch Accuracy (validation) = 0.927697, five\_epoch\_moving\_average = 0.927697  
Epoch: 0021 cost: 0.079925150 Time since last epoch: 6.027366399765015  
epoch Accuracy (validation) = 0.936241, five\_epoch\_moving\_average = 0.930018  
Epoch: 0022 cost: 0.084870398 Time since last epoch: 5.894269943237305  
epoch Accuracy (validation) = 0.931523, five\_epoch\_moving\_average = 0.930834  
Epoch: 0023 cost: 0.086388804 Time since last epoch: 5.9673216342926025  
epoch Accuracy (validation) = 0.938791, five\_epoch\_moving\_average = 0.933027  
Epoch: 0024 cost: 0.079331897 Time since last epoch: 5.944306135177612  
epoch Accuracy (validation) = 0.933053, five\_epoch\_moving\_average = 0.933461  
Epoch: 0025 cost: 0.095811225 Time since last epoch: 5.960317134857178  
epoch Accuracy (validation) = 0.931013, five\_epoch\_moving\_average = 0.934124  
Epoch: 0026 cost: 0.038326349 Time since last epoch: 5.953312635421753  
epoch Accuracy (validation) = 0.930502, five\_epoch\_moving\_average = 0.932976  
Epoch: 0027 cost: 0.063909359 Time since last epoch: 5.984334945678711  
epoch Accuracy (validation) = 0.937899, five\_epoch\_moving\_average = 0.934251  
Epoch: 0028 cost: 0.060087390 Time since last epoch: 5.980331897735596  
epoch Accuracy (validation) = 0.930247, five\_epoch\_moving\_average = 0.932543


Epoch: 0029 cost: 0.068583719 Time since last epoch: 6.011354207992554  
epoch Accuracy (validation) = 0.944147, five\_epoch\_moving\_average = 0.934762  
Epoch: 0030 cost: 0.043039475 Time since last epoch: 6.050382852554321  
New learning rate: 0.001  
epoch Accuracy (validation) = 0.940194, five\_epoch\_moving\_average = 0.936598  
Epoch: 0031 cost: 0.121078894 Time since last epoch: 6.043377637863159  
epoch Accuracy (validation) = 0.940576, five\_epoch\_moving\_average = 0.938613  
Epoch: 0032 cost: 0.044603173 Time since last epoch: 6.070396900177002  
epoch Accuracy (validation) = 0.943254, five\_epoch\_moving\_average = 0.939684  
Epoch: 0033 cost: 0.012664272 Time since last epoch: 6.062391519546509  
epoch Accuracy (validation) = 0.947972, five\_epoch\_moving\_average = 0.943229  
Epoch: 0034 cost: 0.058574189 Time since last epoch: 6.096415758132935  
epoch Accuracy (validation) = 0.944402, five\_epoch\_moving\_average = 0.943280  
Epoch: 0035 cost: 0.041081220 Time since last epoch: 6.079403877258301  
epoch Accuracy (validation) = 0.944274, five\_epoch\_moving\_average = 0.944096  
Epoch: 0036 cost: 0.016299020 Time since last epoch: 6.0313684940338135  
epoch Accuracy (validation) = 0.941724, five\_epoch\_moving\_average = 0.944325  
Epoch: 0037 cost: 0.050895378 Time since last epoch: 6.139447450637817  
epoch Accuracy (validation) = 0.94555, five\_epoch\_moving\_average = 0.944785  
Epoch: 0038 cost: 0.070183635 Time since last epoch: 6.054385662078857  
epoch Accuracy (validation) = 0.942999, five\_epoch\_moving\_average = 0.943790  
Epoch: 0039 cost: 0.187935829 Time since last epoch: 6.092412948608398  
epoch Accuracy (validation) = 0.943892, five\_epoch\_moving\_average = 0.943688  
Epoch: 0040 cost: 0.056463469 Time since last epoch: 6.081405162811279  
epoch Accuracy (validation) = 0.943764, five\_epoch\_moving\_average = 0.943586  
Epoch: 0041 cost: 0.026422201 Time since last epoch: 6.097416639328003  
epoch Accuracy (validation) = 0.95065, five\_epoch\_moving\_average = 0.945371  
Epoch: 0042 cost: 0.024405001 Time since last epoch: 6.1334428787231445  
epoch Accuracy (validation) = 0.948738, five\_epoch\_moving\_average = 0.946009  
Epoch: 0043 cost: 0.037014674 Time since last epoch: 6.145451545715332  
epoch Accuracy (validation) = 0.94708, five\_epoch\_moving\_average = 0.946825  
Epoch: 0044 cost: 0.049144726 Time since last epoch: 6.135444164276123  
epoch Accuracy (validation) = 0.948227, five\_epoch\_moving\_average = 0.947692  
Epoch: 0045 cost: 0.013649072 Time since last epoch: 6.130427837371826  
New learning rate: 0.001  
epoch Accuracy (validation) = 0.94504, five\_epoch\_moving\_average = 0.947947  
Epoch: 0046 cost: 0.039474308 Time since last epoch: 6.152416944503784  
epoch Accuracy (validation) = 0.953073, five\_epoch\_moving\_average = 0.948431  
Epoch: 0047 cost: 0.004283312 Time since last epoch: 6.117392063140869  
epoch Accuracy (validation) = 0.949503, five\_epoch\_moving\_average = 0.948585  
Epoch: 0048 cost: 0.053008489 Time since last epoch: 6.190444469451904  
epoch Accuracy (validation) = 0.943254, five\_epoch\_moving\_average = 0.947819  
Epoch: 0049 cost: 0.013232770 Time since last epoch: 6.1774351596832275  
epoch Accuracy (validation) = 0.945805, five\_epoch\_moving\_average = 0.947335  
Epoch: 0050 cost: 0.014396756 Time since last epoch: 6.202452898025513  
epoch Accuracy (validation) = 0.947717, five\_epoch\_moving\_average = 0.947870  
Epoch: 0051 cost: 0.019178521 Time since last epoch: 6.192445993423462  
epoch Accuracy (validation) = 0.948993, five\_epoch\_moving\_average = 0.947054  
Epoch: 0052 cost: 0.070833355 Time since last epoch: 6.177434921264648  
epoch Accuracy (validation) = 0.949248, five\_epoch\_moving\_average = 0.947003  
Epoch: 0053 cost: 0.005360267 Time since last epoch: 6.196448802947998  
epoch Accuracy (validation) = 0.947845, five\_epoch\_moving\_average = 0.947921  
Epoch: 0054 cost: 0.016064167 Time since last epoch: 6.2395102977752686  
epoch Accuracy (validation) = 0.949248, five\_epoch\_moving\_average = 0.948610  
Epoch: 0055 cost: 0.003761522 Time since last epoch: 6.292487144470215  
epoch Accuracy (validation) = 0.954731, five\_epoch\_moving\_average = 0.950013  
Epoch: 0056 cost: 0.007220398 Time since last epoch: 6.268500328063965

epoch Accuracy (validation) = 0.953073, five\_epoch\_moving\_average = 0.950829  
Epoch: 0057 cost: 0.072295405 Time since last epoch: 6.269501209259033  
epoch Accuracy (validation) = 0.949885, five\_epoch\_moving\_average = 0.950956  
Epoch: 0058 cost: 0.004718235 Time since last epoch: 6.334547996520996  
epoch Accuracy (validation) = 0.94963, five\_epoch\_moving\_average = 0.951313  
Epoch: 0059 cost: 0.110409088 Time since last epoch: 6.3015241622924805  
epoch Accuracy (validation) = 0.954221, five\_epoch\_moving\_average = 0.952308  
Epoch: 0060 cost: 0.083273195 Time since last epoch: 6.325541257858276  
New learning rate: 0.001  
epoch Accuracy (validation) = 0.948865, five\_epoch\_moving\_average = 0.951135  
Epoch: 0061 cost: 0.020478291 Time since last epoch: 6.246484756469727  
epoch Accuracy (validation) = 0.952181, five\_epoch\_moving\_average = 0.950956  
Epoch: 0062 cost: 0.007613658 Time since last epoch: 6.289515018463135  
epoch Accuracy (validation) = 0.95167, five\_epoch\_moving\_average = 0.951313  
Epoch: 0063 cost: 0.055680636 Time since last epoch: 6.239480018615723  
epoch Accuracy (validation) = 0.959194, five\_epoch\_moving\_average = 0.953226  
Epoch: 0064 cost: 0.005754110 Time since last epoch: 6.282510280609131  
epoch Accuracy (validation) = 0.953201, five\_epoch\_moving\_average = 0.953022  
Epoch: 0065 cost: 0.007944511 Time since last epoch: 6.277507305145264  
epoch Accuracy (validation) = 0.9481, five\_epoch\_moving\_average = 0.952869  
Epoch: 0066 cost: 0.013578608 Time since last epoch: 6.311530828475952  
epoch Accuracy (validation) = 0.951033, five\_epoch\_moving\_average = 0.952640  
Epoch: 0067 cost: 0.003912454 Time since last epoch: 6.265498876571655  
epoch Accuracy (validation) = 0.955113, five\_epoch\_moving\_average = 0.953328  
Epoch: 0068 cost: 0.024532700 Time since last epoch: 6.3325464725494385  
epoch Accuracy (validation) = 0.954731, five\_epoch\_moving\_average = 0.952436  
Epoch: 0069 cost: 0.023792870 Time since last epoch: 6.356595993041992  
epoch Accuracy (validation) = 0.95167, five\_epoch\_moving\_average = 0.952130  
Epoch: 0070 cost: 0.003479885 Time since last epoch: 6.378546953201294  
epoch Accuracy (validation) = 0.951415, five\_epoch\_moving\_average = 0.952793  
Epoch: 0071 cost: 0.016517812 Time since last epoch: 6.3875861167907715  
epoch Accuracy (validation) = 0.951288, five\_epoch\_moving\_average = 0.952844  
Epoch: 0072 cost: 0.039033893 Time since last epoch: 6.411603212356567  
epoch Accuracy (validation) = 0.952818, five\_epoch\_moving\_average = 0.952385  
Epoch: 0073 cost: 0.010777423 Time since last epoch: 6.352560758590698  
epoch Accuracy (validation) = 0.956006, five\_epoch\_moving\_average = 0.952640  
Epoch: 0074 cost: 0.093356892 Time since last epoch: 6.407600402832031  
epoch Accuracy (validation) = 0.957664, five\_epoch\_moving\_average = 0.953838  
Epoch: 0075 cost: 0.064308852 Time since last epoch: 6.376578092575073  
New learning rate: 0.001  
epoch Accuracy (validation) = 0.952053, five\_epoch\_moving\_average = 0.953966  
Epoch: 0076 cost: 0.073636942 Time since last epoch: 6.464641094207764  
epoch Accuracy (validation) = 0.957281, five\_epoch\_moving\_average = 0.955164  
Epoch: 0077 cost: 0.115127563 Time since last epoch: 6.392589569091797  
epoch Accuracy (validation) = 0.954221, five\_epoch\_moving\_average = 0.955445  
Epoch: 0078 cost: 0.052979246 Time since last epoch: 6.474648475646973  
epoch Accuracy (validation) = 0.95014, five\_epoch\_moving\_average = 0.954272  
Epoch: 0079 cost: 0.038093597 Time since last epoch: 6.385584115982056  
epoch Accuracy (validation) = 0.955241, five\_epoch\_moving\_average = 0.953787  
Epoch: 0080 cost: 0.041793946 Time since last epoch: 6.468644618988037  
epoch Accuracy (validation) = 0.950013, five\_epoch\_moving\_average = 0.953379  
Epoch: 0081 cost: 0.059610769 Time since last epoch: 6.468644142150879  
epoch Accuracy (validation) = 0.957026, five\_epoch\_moving\_average = 0.953328  
Epoch: 0082 cost: 0.008477603 Time since last epoch: 6.473647832870483  
epoch Accuracy (validation) = 0.953328, five\_epoch\_moving\_average = 0.953150  
Epoch: 0083 cost: 0.070866391 Time since last epoch: 6.517679214477539  
epoch Accuracy (validation) = 0.955751, five\_epoch\_moving\_average = 0.954272

Epoch: 0084 cost: 0.057302065 Time since last epoch: 6.459637403488159  
epoch Accuracy (validation) = 0.952691, five\_epoch\_moving\_average = 0.953762  
Epoch: 0085 cost: 0.067126825 Time since last epoch: 6.455635070800781  
epoch Accuracy (validation) = 0.957409, five\_epoch\_moving\_average = 0.955241  
Epoch: 0086 cost: 0.001766566 Time since last epoch: 6.489658832550049  
epoch Accuracy (validation) = 0.955241, five\_epoch\_moving\_average = 0.954884  
Epoch: 0087 cost: 0.004638392 Time since last epoch: 6.534748554229736  
epoch Accuracy (validation) = 0.955369, five\_epoch\_moving\_average = 0.955292  
Epoch: 0088 cost: 0.005679436 Time since last epoch: 6.519680738449097  
epoch Accuracy (validation) = 0.956899, five\_epoch\_moving\_average = 0.955522  
Epoch: 0089 cost: 0.005729181 Time since last epoch: 6.533705472946167  
epoch Accuracy (validation) = 0.954858, five\_epoch\_moving\_average = 0.955955  
Epoch: 0090 cost: 0.010331319 Time since last epoch: 6.542698621749878  
New learning rate: 0.001  
epoch Accuracy (validation) = 0.954858, five\_epoch\_moving\_average = 0.955445  
Epoch: 0091 cost: 0.007515794 Time since last epoch: 6.452662229537964  
epoch Accuracy (validation) = 0.952181, five\_epoch\_moving\_average = 0.954833  
Epoch: 0092 cost: 0.005800207 Time since last epoch: 6.536692142486572  
epoch Accuracy (validation) = 0.954348, five\_epoch\_moving\_average = 0.954629  
Epoch: 0093 cost: 0.062640086 Time since last epoch: 6.556677341461182  
epoch Accuracy (validation) = 0.958174, five\_epoch\_moving\_average = 0.954884  
Epoch: 0094 cost: 0.034962095 Time since last epoch: 6.508673429489136  
epoch Accuracy (validation) = 0.956899, five\_epoch\_moving\_average = 0.955292  
Epoch: 0095 cost: 0.003305486 Time since last epoch: 6.5687150955200195  
epoch Accuracy (validation) = 0.956389, five\_epoch\_moving\_average = 0.955598  
Epoch: 0096 cost: 0.046248607 Time since last epoch: 6.661783933639526  
epoch Accuracy (validation) = 0.954986, five\_epoch\_moving\_average = 0.956159  
Epoch: 0097 cost: 0.003345174 Time since last epoch: 6.62578821182251  
epoch Accuracy (validation) = 0.957664, five\_epoch\_moving\_average = 0.956822  
Epoch: 0098 cost: 0.015654625 Time since last epoch: 6.5707175731658936  
epoch Accuracy (validation) = 0.957026, five\_epoch\_moving\_average = 0.956593  
Epoch: 0099 cost: 0.016251313 Time since last epoch: 6.519681215286255  
epoch Accuracy (validation) = 0.953838, five\_epoch\_moving\_average = 0.955981  
Epoch: 0100 cost: 0.051637076 Time since last epoch: 6.615718126296997  
epoch Accuracy (validation) = 0.95167, five\_epoch\_moving\_average = 0.955037  
Epoch: 0101 cost: 0.016700024 Time since last epoch: 6.605775356292725  
epoch Accuracy (validation) = 0.958174, five\_epoch\_moving\_average = 0.955675  
Epoch: 0102 cost: 0.024302680 Time since last epoch: 6.539692640304565  
epoch Accuracy (validation) = 0.959959, five\_epoch\_moving\_average = 0.956134  
Epoch: 0103 cost: 0.007559350 Time since last epoch: 6.637734889984131  
epoch Accuracy (validation) = 0.950268, five\_epoch\_moving\_average = 0.954782  
Epoch: 0104 cost: 0.006651680 Time since last epoch: 6.662298917770386  
epoch Accuracy (validation) = 0.956516, five\_epoch\_moving\_average = 0.955317  
Epoch: 0105 cost: 0.016676918 Time since last epoch: 6.683793544769287  
New learning rate: 0.001  
epoch Accuracy (validation) = 0.960214, five\_epoch\_moving\_average = 0.957026  
Epoch: 0106 cost: 0.003506209 Time since last epoch: 6.6568920612335205  
epoch Accuracy (validation) = 0.958301, five\_epoch\_moving\_average = 0.957052  
Epoch: 0107 cost: 0.011197917 Time since last epoch: 6.675769329071045  
epoch Accuracy (validation) = 0.959194, five\_epoch\_moving\_average = 0.956899  
Epoch: 0108 cost: 0.007262496 Time since last epoch: 6.699810981750488  
epoch Accuracy (validation) = 0.957664, five\_epoch\_moving\_average = 0.958378  
Epoch: 0109 cost: 0.008470118 Time since last epoch: 6.738868474960327  
epoch Accuracy (validation) = 0.956006, five\_epoch\_moving\_average = 0.958276  
Epoch: 0110 cost: 0.002060802 Time since last epoch: 6.6337316036224365  
epoch Accuracy (validation) = 0.957281, five\_epoch\_moving\_average = 0.957689  
Epoch: 0111 cost: 0.108816624 Time since last epoch: 6.636765241622925



epoch Accuracy (validation) = 0.956771, five\_epoch\_moving\_average = 0.957383  
Epoch: 0112 cost: 0.026132941 Time since last epoch: 6.74183988571167  
epoch Accuracy (validation) = 0.959959, five\_epoch\_moving\_average = 0.957536  
Epoch: 0113 cost: 0.093736030 Time since last epoch: 6.73883843421936  
epoch Accuracy (validation) = 0.958684, five\_epoch\_moving\_average = 0.957740  
Epoch: 0114 cost: 0.001585683 Time since last epoch: 6.722826719284058  
epoch Accuracy (validation) = 0.958684, five\_epoch\_moving\_average = 0.958276  
Epoch: 0115 cost: 0.002818441 Time since last epoch: 6.809889078140259  
epoch Accuracy (validation) = 0.961107, five\_epoch\_moving\_average = 0.959041  
Epoch: 0116 cost: 0.061915379 Time since last epoch: 6.718854188919067  
epoch Accuracy (validation) = 0.961999, five\_epoch\_moving\_average = 0.960087  
Optimization Finished!



## 4.4 Test Model

```
In [19]: # Test model
correct_prediction = tf.equal(tf.argmax(pred, 1), tf.argmax(y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))

with sess.as_default():
    print("Accuracy (test):", accuracy.eval({x_unflattened: X_test, y_rawlabels: y_test}))

# Print parameters for reference
print("\nParameters:")
print("Learning rate (initial): ", initial_learning_rate)
print("Anneal learning rate every ", anneal_mod_frequency, " epochs by ", 1 - annealing_rate)
print("Learning rate (final): ", learning_rate)
print("Training epochs: ", training_epochs)
print("Batch size: ", batch_size)
print("Dropout (conv): ", dropout_conv)
print("Dropout (fc): ", dropout_fc)
print("Padding: ", padding)
print("weights_mean: ", weights_mean)
print("weights_stddev: ", weights_stddev)
print("biases_mean: ", biases_mean)
```

Accuracy (test): 0.793508

Parameters:  
 Learning rate (initial): 0.001  
 Anneal learning rate every 15 epochs by 0  
 Learning rate (final): 0.001  
 Training epochs: 150  
 Batch size: 100  
 Dropout (conv): 0.9  
 Dropout (fc): 0.9  
 Padding: VALID  
 weights\_mean: 0.0  
 weights\_stddev: 0.1  
 biases\_mean: 0.0

## Question 1

*Describe how you preprocessed the data. Why did you choose that technique?*

**Answer:**

See Section 2.1 for Preprocessing steps

1. As image was 32x32 no Padding done
2. avoid overfitting by features randomization
3. data normalized to avoid high variance and improve classifier performance
4. labels one hot encoded
5. Added Extra data for low frequency Sign classes

## Question 2

*Describe how you set up the training, validation and testing data for your model. **Optional:** If you generated additional data, how did you generate the data? Why did you generate the data? What are the differences in the new dataset (with generated data) from the original dataset?*

**Answer:**

See Section 2.1 for Splitting Data

1. Training and test data were already separated (downloaded pickled files train.p and test.p).
2. shuffled the training data because they were arranged in ascending order by label. If I don't shuffle the training data, the first series of batches will all be the first type of sign followed by the second type and so on. This will distort the learning process.
3. I further split the training data into test and validation sets so the model wouldn't be cheating when we optimised it.
4. Improvement: Generate additional data. The number of examples in the training data for each class is uneven, so the model may be biased towards predicting an unknown sign belongs to a class where there is abundant training data since we are minimising the training loss.

## Question 3

*What does your final architecture look like? (Type of model, layers, sizes, connectivity, etc.) For reference on how to build a deep neural network using TensorFlow, see [Deep Neural Network in TensorFlow](https://classroom.udacity.com/nanodegrees/nd013/parts/fbf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/b516a270-8600-4f93-a0a3-20dfeabe5da6/concepts/83a3a2a2-a9bd-4b7b-95b0-eb924ab14432) (<https://classroom.udacity.com/nanodegrees/nd013/parts/fbf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/b516a270-8600-4f93-a0a3-20dfeabe5da6/concepts/83a3a2a2-a9bd-4b7b-95b0-eb924ab14432>) from the classroom.*

**Answer:**

See 3.3 for Model / Network Definition

- 3-layer Convolutional Neural Network.
- It consists of one convolution layer (feature extraction) followed by two fully connected layers (ReLU activation) and a single fully connected linear classifier.

1 : Convolution layer Input: (32, 32, 3)

Output: (5, 5, 32) 'VALID' padding Filters: 32 Stride: 3 ReLU Activation 2D Max Pooling (down-sampling) layer Dropout: 0.9

2 : Reshape Layer Input: (5, 5, 32) Output: 800

3 : Fully connected layer

Input: 800 Output: 512 ReLU Activation Dropout: 0.9

4 : Fully connected layer

Input: 512 Output: 128 ReLU Activation Dropout: 0.9

5 : output layer

Input: 128 Output: 43

- The network uses full colour information (all three channels) and normalised data.

## Question 4

*How did you train your model? (Type of optimizer, batch size, epochs, hyperparameters, etc.)*

See Section 4.3

- Type of optimiser: AdamOptimizer
- Batch size: 100
- Training Epochs: 105(with ELU)
- Learning rate: 0.001

Network Parameters:

- Dropout (conv layer): 0.9
- Dropout (fully connected layers): 0.9
- Padding: VALID

tf.train.AdamOptimizer ref <http://stats.stackexchange.com/questions/184448/difference-between-gradientdescentoptimizer-and-adamoptimizer-tensorflow>  
(<http://stats.stackexchange.com/questions/184448/difference-between-gradientdescentoptimizer-and-adamoptimizer-tensorflow>)

•Main advantage of Adam over the simple tf.train.GradientDescentOptimizer: Uses moving averages of the parameters (momentum) -> enables Adam to use a larger effective step size, and the algorithm will converge to this step size without fine tuning. A simple tf.train.GradientDescentOptimizer would require more hyperparameter tuning before it would converge as quickly.

•Disadvantage: Adam requires more computation to be performed for each parameter in each training step (to maintain the moving averages and variance, and calculate the scaled gradient) and more state to be retained for each parameter (approximately tripling the size of the model to store the average and variance for each parameter).

## Question 5

*What approach did you take in coming up with a solution to this problem? It may have been a process of trial and error, in which case, outline the steps you took to get to the final solution and why you chose those steps.*

*Perhaps your solution involved an already well known implementation or architecture. In this case, discuss why you think this is suitable for the current problem.*

## **Answer:**

### **1. First attempt: building a minimum viable model and debugging**

- I wanted to get a working model first. I started with a basic multilayer perceptron which I adapted from TensorFlow-Examples. I trained it for 15 epochs, which had an accuracy of 6% on the training and test sets. I then trained a two-layer convolutional neural network for 15 epochs which had an accuracy of 5-6% on the training and test sets.
  - The accuracy was lower than I expected and the cost seemed high (of order  $10^6$  in the first epoch,  $10^5$  in the second and third and in the hundreds in the tenth epoch), so I adjusted parameters hoping to improve it before training for longer.
  - The cost reduced significantly (to single digits by the second epoch as opposed to order  $10^5$ ) after I added a small positive bias to the initial weights and biases. Strangely, the accuracy did not increase, but remained at 5-6%. The cost did not decrease significantly over the next 10 epochs either.
  - I went on Slack to see what results people were getting to get a feel for how wrong I was. I saw that people often trained their networks for hundreds of epochs so I thought it would be good to train my network for e.g. 100 epochs.
- I rewrote the multilayer perceptron in a Python Script and it worked fine, returning an accuracy of over 70% accuracy within 2 epochs.

### **2. Improvements to the model**

- I then added a convolution layer before the two fully connected layers and the output layer.
- This new model produced a validation accuracy of above 90% after 15 epochs (parameters not tuned), which was higher than that for the two-layer feedforward network. So I chose this model with a convolution layer.

### **3. Tuning Parameters**

- I altered the model code (replaced hard-coded numbers with variables) so I could tweak parameters easily.
- I tested models with different values or settings for
  - dropout for the fully connected layers,
  - dropout for the convolution layer,
  - padding (SAME vs VALID),
  - weight and bias initialisation
  - maxpool vs no max pool
- I used Keras to implement comparisons so I could get full figures on training and validation loss and accuracy easily.
- I stopped when the model reached a validation accuracy of over 95% within 100 epochs.
  - This figure is strange because my models implemented in Keras reach validation accuracy of over 99% within 15 epochs.

---

## Step 3: Test a Model on New Images

Take several pictures of traffic signs that you find on the web or around you (at least five), and run them through your classifier on your computer to produce example results. The classifier might not recognize some local signs but it could prove interesting nonetheless.

You may find `signnames.csv` useful as it contains mappings from the class id (integer) to the actual sign name.

### Implementation

Use the code cell (or multiple code cells, if necessary) to implement the first step of your project. Once you have completed your implementation and are satisfied with the results, be sure to thoroughly answer the questions that follow.

```

In [20]: def read_image_and_print_dims(image_path):
            image = mpimg.imread(image_path)
            print('This image is:', type(image), 'with dimensions:', image.shape)
            plt.imshow(image)
            return image

        def process_image_file(name):
            image = cv2.imread(name)
            image = cv2.resize(image,(32,32))
            image = cv2.cvtColor(image,cv2.COLOR_BGR2RGB)
            image = image/255.-.5
            return image

        def plot_image_grid(n_row,n_col,X):
            plt.figure(figsize = (8,6))
            gs1 = gridspec.GridSpec(n_row,n_col)
            gs1.update(wspace=0.01, hspace=0.02) # set the spacing between axes.

            for i in range(n_row*n_col):
                # i = i + 1 # grid spec indexes from 0
                ax1 = plt.subplot(gs1[i])
                plt.axis('on')
                ax1.set_xticklabels([])
                ax1.set_yticklabels([])
                ax1.set_aspect('equal')
                #plt.subplot(4,11,i+1)
                ind_plot = i
                plt.imshow(X[ind_plot])
                plt.axis('off')
            plt.show()

        def predict(img):
            classification = sess.run(tf.argmax(pred, 1), feed_dict={x_unflattened: [img]})
            print('NN predicted', classification[0])

```



```

plt.imshow(sign_dict[classification[0]])
plt.show()

def top_5_predictions(img):
    #Return model's top five choices for what traffic sign this image is and its confidence in its predictions.
    top_five_certainities = sess.run(tf.nn.top_k(tf.nn.softmax(pred), k=5), feed_dict={x_unflattened: [img]})
    print("Top five: ", top_five_certainities)
    plot_certainty_arrays(top_five_certainities[0][0], top_five_certainities[1][0])
    return top_five_certainities

def plot_certainty_arrays(probabilities, labels):
    # Plot model's probabilities (y) and traffic sign labels (x) in a bar chart.
    y_pos = np.arange(len(labels))
    plt.bar(y_pos, probabilities, align='center', alpha=0.5)
    plt.xticks(y_pos, labels)
    plt.ylabel('Probability')
    plt.xlabel('Traffic sign')
    plt.title('Model\'s certainty of its predictions')
    plt.show()
    print("Traffic Sign Key")
    for label in labels:
        print(label, ": ", data_pd.loc[label]['SignName'])

```

## Question 6

Choose five candidate images of traffic signs and provide them in the report. Are there any particular qualities of the image(s) that might make classification difficult? It could be helpful to plot the images in the notebook.

**Answer:**

```

In [21]: newdata = [process_image_file("./new_signs/"+name) for name in os.listdir("./new_signs/")]
          namenewdata = [name for name in os.listdir("./new_signs/")]
          newdata = np.array(newdata, dtype = np.float32)
          plot_image_grid(9,5,newdata+.5)

```



## Question 7


*Is your model able to perform equally well on captured pictures when compared to testing on the dataset? The simplest way to do this check the accuracy of the predictions. For example, if the model predicted 1 out of 5 signs correctly, it's 20% accurate.*

**NOTE:** *You could check the accuracy manually by using `signnames.csv` (same directory). This file has a mapping from the class id (0-42) to the corresponding sign name. So, you could take the class id the model outputs, lookup the name in `signnames.csv` and see if it matches the sign from the image.*


**Answer:**

```
In [22]: for i in range(len(namenewdata)):
          print('\n\n\n\n')
          print(namenewdata[i])
          plt.imshow(newdata[i]+.5)
          plt.show()
          predict(newdata[i])
```


40mph.png




NN predicted 25




50mph.png




NN predicted 25




8pcGrade.png




NN predicted 5




bikelane.png




NN predicted 25




circle.png




NN predicted 12




CrossRoad.png



NN predicted 20




DonotEnter.png




NN predicted 17







DoNotPass.png




NN predicted 20




EZPass.png




NN predicted 13




german\_sign.png




NN predicted 30




HCparking.png




NN predicted 33




HighXWinds.png




NN predicted 26




Hospital.png




NN predicted 12




HOW.png




NN predicted 33



i1.png




NN predicted 35




i22.png



NN predicted 38



japanese\_sign\_resized.png




NN predicted 14




keepleft.png







NN predicted 13




keepright.png




NN predicted 13




noleft.png




NN predicted 2




noleftUturn.png




NN predicted 9




NoParking.png




NN predicted 33




noright.png




NN predicted 33




noRoller.png




NN predicted 9




nostraight.png




NN predicted 1




noUturn.png




NN predicted 35




noWalking.png




NN predicted 1



paidParking.png




NN predicted 33







parking2hr.png




NN predicted 25




RailRoad.png




NN predicted 12




RightCrossing.png




NN predicted 20




roadwork.png




NN predicted 12




shark\_sign.png




NN predicted 12




speed\_limit\_stop.png




NN predicted 12




stop.png




NN predicted 1




StopAhead.png




NN predicted 12




TowAway.png




NN predicted 22




two\_way\_sign.png



NN predicted 0




US50.png




NN predicted 21







US6.png




NN predicted 11




windyRight.png




NN predicted 12




workzone.png




NN predicted 12




YCrossing.png




NN predicted 12




yeild.png




NN predicted 13



yield\_pedestrian.png



NN predicted 12



No, it does not perform equally well on captured images. It has a performance of 0% accuracy on captured images as opposed to 79% on the test set. •The images not included in the dataset are not exactly the same road signs so there is additional difficulty because the model needs to generalise well to classify these new signs correctly. The •Some road signs such as the shark sign may not even be included in the 43 categories. •The images are also processed (e.g. cropped) differently.

It seems that the model is classifying 'unknown signs' as Roundabout Mandatory signs.

## Question 8

Use the model's softmax probabilities to visualize the **certainty** of its predictions, `tf.nn.top_k` ([https://www.tensorflow.org/versions/r0.12/api\\_docs/python/nn.html#top\\_k](https://www.tensorflow.org/versions/r0.12/api_docs/python/nn.html#top_k)) could prove helpful here. Which predictions is the model certain of? Uncertain? If the model was incorrect in its initial prediction, does the correct prediction appear in the top k? (k should be 5 at most)

`tf.nn.top_k` will return the values and indices (class ids) of the top k predictions. So if k=3, for each sign, it'll return the 3 largest probabilities (out of a possible 43) and the corresponding class ids.

Take this numpy array as an example:

```
# (5, 6) array
a = np.array([[ 0.24879643,  0.07032244,  0.12641572,  0.34763842,  0.07893497,
                0.12789202],
              [ 0.28086119,  0.27569815,  0.08594638,  0.0178669 ,  0.18063401,
                0.15899337],
              [ 0.26076848,  0.23664738,  0.08020603,  0.07001922,  0.1134371 ,
                0.23892179],
              [ 0.11943333,  0.29198961,  0.02605103,  0.26234032,  0.1351348 ,
                0.16505091],
              [ 0.09561176,  0.34396535,  0.0643941 ,  0.16240774,  0.24206137,
                0.09155967]])
```

Running it through `sess.run(tf.nn.top_k(tf.constant(a), k=3))` produces:

```
TopKV2(values=array([[ 0.34763842,  0.24879643,  0.12789202],
                    [ 0.28086119,  0.27569815,  0.18063401],
                    [ 0.26076848,  0.23892179,  0.23664738],
                    [ 0.29198961,  0.26234032,  0.16505091],
                    [ 0.34396535,  0.24206137,  0.16240774]]), indices=array([[3, 0, 5],
                                     [0, 1, 4],
                                     [0, 5, 1],
                                     [1, 3, 5],
                                     [1, 4, 3]], dtype=int32))
```

Looking just at the first row we get `[ 0.34763842, 0.24879643, 0.12789202]`, you can confirm these are the 3 largest probabilities in `a`. You'll also notice `[3, 0, 5]` are the corresponding indices.

### Answer:

- The model is certain of all of its predictions even though some are wrong.
- The model also predicts different outcomes confidently for the two times I ran the predictions on each sign.


These are both strange outcomes.

**Note:** Once you have completed all of the code implementations and successfully answered each question above, you may finalize your work by exporting the iPython Notebook as an HTML document. You can do this by using the menu above and navigating to \n", "**File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.


```
In [23]: for i in range(len(namenewdata)):
          print('\n\n\n\n')
          print(namenewdata[i])
          plt.imshow(newdata[i]+.5)
          plt.show()
          top_5_predictions(newdata[i])
```



40mph.png




Top five: TopKV2(values=array([[ 9.99965072e-01, 3.48573485e-05, 9.53655004e-08, 5.30022311e-11, 4.23061922e-14]], dtype=float32), indices=array([[25, 28, 13, 14, 12]]))




Traffic Sign Key

25 : Road work  
28 : Children crossing  
13 : Yield  
14 : Stop  
12 : Priority road

50mph.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 1.46745028e-29, 1.35838156e-29, 1.21232870e-30, 4.28406213e-32]], dtype=float32), indices=array([[25, 33, 13, 40, 19]]))




Traffic Sign Key  
 25 : Road work  
 33 : Turn right ahead  
 13 : Yield  
 40 : Roundabout mandatory  
 19 : Dangerous curve to the left

8pcGrade.png



Top five: TopKV2(values=array([[ 9.34373081e-01, 6.01442643e-02, 5.23917377e-03, 2.43519113e-04, 1.50220853e-14]], dtype=float32), indices=array([[12, 20, 5, 3, 42]]))



Traffic Sign Key

12 : Priority road


20 : Dangerous curve to the right

5 : Speed limit (80km/h)


3 : Speed limit (60km/h)

42 : End of no passing by vehicles over 3.5 metric tons

bikelane.png




Top five: TopKV2(values=array([[ 9.99993324e-01, 6.67676613e-06, 1.72000980e-09, 2.67864536e-10, 1.98939309e-10]], dtype=float32), indices=array([[12, 26, 28, 25, 20]]))




Traffic Sign Key  
 12 : Priority road  
 26 : Traffic signals  
 28 : Children crossing  
 25 : Road work  
 20 : Dangerous curve to the right

circle.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 1.11131870e-33, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00], dtype=float32), indices=array([[12, 8, 0, 1, 2]]))




Traffic Sign Key  
 12 : Priority road  
 8 : Speed limit (120km/h)  
 0 : Speed limit (20km/h)  
 1 : Speed limit (30km/h)  
 2 : Speed limit (50km/h)

CrossRoad.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 6.52831078e-15, 6.28363006e-15, 7.22331169e-16, 5.61500953e-18]], dtype=float32), indices=array([[12, 25, 20, 30, 31]]))




Traffic Sign Key  
 12 : Priority road  
 25 : Road work  
 20 : Dangerous curve to the right  
 30 : Beware of ice/snow  
 31 : Wild animals crossing

DonotEnter.png




Top five: TopKV2(values=array([[ 9.99999523e-01, 4.68118913e-07, 7.32648371e-13, 1.17335677e-17, 8.93368379e-18]], dtype=float32), indices=array([[17, 25, 1, 14, 2]]))




Traffic Sign Key

17 : No entry  
 25 : Road work  
 1 : Speed limit (30km/h)  
 14 : Stop  
 2 : Speed limit (50km/h)

DoNotPass.png



Top five: TopKV2(values=array([[ 9.99998689e-01, 8.35011463e-07, 4.23040120e-07, 5.35495148e-10, 7.56928201e-11]], dtype=float32), indices=array([[20, 3, 10, 21, 31]]))




#### Traffic Sign Key


20 : Dangerous curve to the right  
 3 : Speed limit (60km/h)  
 10 : No passing for vehicles over 3.5 metric tons  
 21 : Double curve  
 31 : Wild animals crossing

EZPass.png






Top five: TopKV2(values=array([[ 1.00000000e+00, 3.34847494e-20, 3.01740457e-22, 2.34127782e-22, 5.46599749e-24]], dtype=float32), indices=array([[13, 5, 12, 38, 25]]))




Traffic Sign Key  
 13 : Yield  
 5 : Speed limit (80km/h)  
 12 : Priority road  
 38 : Keep right  
 25 : Road work

german\_sign.png



Top five: TopKV2(values=array([[ 9.59407389e-01, 3.94117013e-02, 1.05212280e-03, 9.47257286e-05, 1.95071989e-05]], dtype=float32), indices=array([[28, 30, 39, 41, 16]]))



Traffic Sign Key

28 : Children crossing


30 : Beware of ice/snow

39 : Keep left


41 : End of no passing

16 : Vehicles over 3.5 metric tons prohibited

HCparking.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 1.62666736e-09, 1.30975786e-11, 1.23612373e-12, 6.92642399e-17]], dtype=float32), indices=array([[33, 35, 20, 40, 23]]))




Traffic Sign Key  
 33 : Turn right ahead  
 35 : Ahead only  
 20 : Dangerous curve to the right  
 40 : Roundabout mandatory  
 23 : Slippery road

HighXWinds.png




Top five: TopKV2(values=array([[ 9.97674763e-01, 1.59696769e-03, 7.22067547e-04, 6.17395062e-06, 3.94209491e-12]], dtype=float32), indices=array([[12, 26, 25, 5, 22]]))




Traffic Sign Key  
 12 : Priority road  
 26 : Traffic signals  
 25 : Road work  
 5 : Speed limit (80km/h)  
 22 : Bumpy road

Hospital.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 3.00917657e-13, 7.56618109e-24, 1.29688733e-24, 1.29791112e-25]], dtype=float32), indices=array([[12, 17, 23, 5, 9]]))




Traffic Sign Key  
 12 : Priority road  
 17 : No entry  
 23 : Slippery road  
 5 : Speed limit (80km/h)  
 9 : No passing

HOW.png




Top five: TopKV2(values=array([[ 7.71307647e-01, 2.28692353e-01, 6.40906439e-09, 3.23991187e-13, 1.66075717e-14]], dtype=float32), indices=array([[33, 25, 22, 13, 12]]))




Traffic Sign Key  
 33 : Turn right ahead  
 25 : Road work  
 22 : Bumpy road  
 13 : Yield  
 12 : Priority road

i1.png



Top five: `TopKV2(values=array([[ 1.00000000e+00, 4.19434503e-18, 1.08677887e-29, 6.71084346e-30, 1.48594288e-30]], dtype=float32), indices=array([[35, 33, 40, 37, 36]]))`




Traffic Sign Key  
 35 : Ahead only  
 33 : Turn right ahead  
 40 : Roundabout mandatory  
 37 : Go straight or left  
 36 : Go straight or right

i22.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 6.12729160e-13, 9.18513549e-15, 6.62731006e-19, 4.33373360e-25]], dtype=float32), indices=array([[38, 41, 13, 40, 3]]))




Traffic Sign Key  
 38 : Keep right  
 41 : End of no passing  
 13 : Yield  
 40 : Roundabout mandatory  
 3 : Speed limit (60km/h)

japanese\_sign\_resized.png





Top five: TopKV2(values=array([[ 1.00000000e+00, 5.30353143e-23, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00]], dtype=float32), indices=array([[17, 14, 0, 1, 2]]))



Traffic Sign Key

17 : No entry


14 : Stop

0 : Speed limit (20km/h)


1 : Speed limit (30km/h)

2 : Speed limit (50km/h)

keepleft.png



Top five: TopKV2(values=array([[ 9.99683142e-01, 3.16423189e-04, 4.58910677e-07, 5.36083290e-14, 8.27646354e-20]], dtype=float32), indices=array([[13, 14, 15, 10, 25]]))



Traffic Sign Key

13 : Yield


14 : Stop

15 : No vehicles


10 : No passing for vehicles over 3.5 metric tons

25 : Road work

keepright.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 6.64526323e-10, 3.46017215e-10, 3.10718324e-15, 1.88433491e-16]], dtype=float32), indices=array([[13, 9, 12, 23, 25]]))




Traffic Sign Key  
 13 : Yield  
 9 : No passing  
 12 : Priority road  
 23 : Slippery road  
 25 : Road work

noleft.png




Top five: TopKV2(values=array([[ 9.99995589e-01, 4.41855764e-06, 6.24935936e-11, 8.73254207e-15, 1.22476697e-15]], dtype=float32), indices=array([[ 2, 33, 17, 9, 4]]))




Traffic Sign Key  
 2 : Speed limit (50km/h)  
 33 : Turn right ahead  
 17 : No entry  
 9 : No passing  
 4 : Speed limit (70km/h)

noleftUturn.png



Top five: TopKV2(values=array([[ 9.57131743e-01, 3.91714312e-02, 3.09280888e-03, 4.61697055e-04, 1.42223202e-04]], dtype=float32), indices=array([[ 2, 9, 17, 16, 0]]))



Traffic Sign Key

2 : Speed limit (50km/h)


9 : No passing

17 : No entry


16 : Vehicles over 3.5 metric tons prohibited

0 : Speed limit (20km/h)

NoParking.png




Top five: TopKV2(values=array([[ 9.99999642e-01, 3.51883159e-07, 1.33628820e-11, 8.08097946e-12, 6.12855469e-12]], dtype=float32), indices=array([[10, 11, 35, 33, 13]]))




Traffic Sign Key

- 10 : No passing for vehicles over 3.5 metric tons
- 11 : Right-of-way at the next intersection
- 35 : Ahead only
- 33 : Turn right ahead
- 13 : Yield

noright.png




Top five: TopKV2(values=array([[ 9.99990463e-01, 6.16097032e-06, 3.39195935e-06, 4.63099292e-09, 8.33453098e-11]], dtype=float32), indices=array([[14, 25, 4, 18, 17]]))




Traffic Sign Key

14 : Stop  
 25 : Road work  
 4 : Speed limit (70km/h)  
 18 : General caution  
 17 : No entry

noRoller.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 1.76973664e-16, 2.10716522e-20, 3.50180884e-32, 4.68891931e-34]], dtype=float32), indices=array([[25, 9, 1, 41, 14]]))




Traffic Sign Key  
 25 : Road work  
 9 : No passing  
 1 : Speed limit (30km/h)  
 41 : End of no passing  
 14 : Stop

nostraight.png





Top five: TopKV2(values=array([[ 4.55018729e-01, 3.59127194e-01, 1.85007110e-01, 7.56068039e-04, 4.96152534e-05]], dtype=float32), indices=array([[24, 1, 26, 18, 14]]))



Traffic Sign Key

24 : Road narrows on the right


1 : Speed limit (30km/h)

26 : Traffic signals


18 : General caution

14 : Stop

noUturn.png




Top five: TopKV2(values=array([[ 9.99811232e-01, 1.67322883e-04, 2.08163110e-05, 6.42379177e-07, 6.93807500e-09]], dtype=float32), indices=array([[16, 10, 9, 12, 42]]))




#### Traffic Sign Key

- 16 : Vehicles over 3.5 metric tons prohibited
- 10 : No passing for vehicles over 3.5 metric tons
- 9 : No passing
- 12 : Priority road
- 42 : End of no passing by vehicles over 3.5 metric tons

noWalking.png




Top five: TopKV2(values=array([[ 9.99984145e-01, 1.53274432e-05, 4.44997141e-07, 4.76872763e-08, 2.59136844e-11]], dtype=float32), indices=array([[ 1, 3, 2, 5, 21]]))




Traffic Sign Key

1 : Speed limit (30km/h)  
 3 : Speed limit (60km/h)  
 2 : Speed limit (50km/h)  
 5 : Speed limit (80km/h)  
 21 : Double curve

paidParking.png




Top five: TopKV2(values=array([[ 9.99614239e-01, 3.81746097e-04, 3.84074065e-06, 9.72145244e-08, 7.92053445e-09]], dtype=float32), indices=array([[18, 33, 37, 26, 17]]))




Traffic Sign Key  
 18 : General caution  
 33 : Turn right ahead  
 37 : Go straight or left  
 26 : Traffic signals  
 17 : No entry

parking2hr.png




Top five: TopKV2(values=array([[ 9.99770820e-01, 2.25819371e-04, 3.05190792e-06, 2.56343213e-07, 7.55058245e-16]], dtype=float32), indices=array([[13, 18, 38, 25, 37]]))




Traffic Sign Key  
 13 : Yield  
 18 : General caution  
 38 : Keep right  
 25 : Road work  
 37 : Go straight or left

RailRoad.png



Top five: `TopKV2(values=array([[ 9.99515653e-01, 4.32657776e-04, 5.11681137e-05, 5.79016614e-07, 1.90898430e-09]], dtype=float32), indices=array([[12, 20, 11, 25, 30]]))`



Traffic Sign Key

12 : Priority road


20 : Dangerous curve to the right

11 : Right-of-way at the next intersection


25 : Road work

30 : Beware of ice/snow

RightCrossing.png



Top five: TopKV2(values=array([[ 1.00000000e+00, 8.59899988e-22, 2.71822030e-22, 1.24089544e-27, 3.18170511e-30]], dtype=float32), indices=array([[20, 26, 12, 8, 42]]))



#### Traffic Sign Key

20 : Dangerous curve to the right


26 : Traffic signals

12 : Priority road


8 : Speed limit (120km/h)

42 : End of no passing by vehicles over 3.5 metric tons

roadwork.png




Top five: TopKV2(values=array([[ 9.97832477e-01, 2.16756505e-03, 1.66943620e-23, 3.95174151e-25, 1.41113320e-27]], dtype=float32), indices=array([[12, 25, 14, 29, 18]]))




Traffic Sign Key  
 12 : Priority road  
 25 : Road work  
 14 : Stop  
 29 : Bicycles crossing  
 18 : General caution

shark\_sign.png






Top five: TopKV2(values=array([[ 1.00000000e+00, 3.56363076e-20, 4.22397344e-31, 8.22733639e-33, 5.29473108e-33]], dtype=float32), indices=array([[12, 25, 17, 30, 20]]))




Traffic Sign Key  
 12 : Priority road  
 25 : Road work  
 17 : No entry  
 30 : Beware of ice/snow  
 20 : Dangerous curve to the right

speed\_limit\_stop.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 5.79819976e-26, 2.64420888e-34, 1.73645782e-34, 0.00000000e+00]], dtype=float32), indices=array([[12, 26, 20, 22, 0]]))




Traffic Sign Key  
 12 : Priority road  
 26 : Traffic signals  
 20 : Dangerous curve to the right  
 22 : Bumpy road  
 0 : Speed limit (20km/h)

stop.png




Top five: TopKV2(values=array([[ 8.73517990e-01, 1.26350164e-01, 1.25370047e-04, 6.48536616e-06, 1.15652359e-08]], dtype=float32), indices=array([[25, 24, 1, 14, 13]]))




Traffic Sign Key  
 25 : Road work  
 24 : Road narrows on the right  
 1 : Speed limit (30km/h)  
 14 : Stop  
 13 : Yield

StopAhead.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 5.42289935e-10, 1.05783150e-30, 0.00000000e+00, 0.00000000e+00]], dtype=float32), indices=array([[12, 14, 25, 0, 1]]))




Traffic Sign Key  
 12 : Priority road  
 14 : Stop  
 25 : Road work  
 0 : Speed limit (20km/h)  
 1 : Speed limit (30km/h)

TowAway.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 8.72576265e-28, 5.75557000e-37, 0.00000000e+00, 0.00000000e+00]], dtype=float32), indices=array([[12, 22, 13, 0, 1]]))




Traffic Sign Key  
 12 : Priority road  
 22 : Bumpy road  
 13 : Yield  
 0 : Speed limit (20km/h)  
 1 : Speed limit (30km/h)

two\_way\_sign.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 1.29474698e-09, 2.46179077e-10, 1.58991700e-12, 5.65064525e-22]], dtype=float32), indices=array([[20, 8, 26, 0, 12]]))




Traffic Sign Key  
 20 : Dangerous curve to the right  
 8 : Speed limit (120km/h)  
 26 : Traffic signals  
 0 : Speed limit (20km/h)  
 12 : Priority road

US50.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 1.77870145e-08, 9.34713529e-10, 1.83235033e-13, 7.02551962e-14]], dtype=float32), indices=array([[30, 25, 24, 23, 13]]))




Traffic Sign Key  
 30 : Beware of ice/snow  
 25 : Road work  
 24 : Road narrows on the right  
 23 : Slippery road  
 13 : Yield

US6.png



Top five: TopKV2(values=array([[ 1.00000000e+00, 3.07357118e-09, 1.45803882e-19, 3.46409604e-28, 1.43348339e-30]], dtype=float32), indices=array([[11, 19, 30, 25, 0]]))




#### Traffic Sign Key


- 11 : Right-of-way at the next intersection
- 19 : Dangerous curve to the left
- 30 : Beware of ice/snow
- 25 : Road work
- 0 : Speed limit (20km/h)

windyRight.png






Top five: TopKV2(values=array([[ 9.42513764e-01, 5.74862510e-02, 5.99843388e-17, 2.31940568e-24, 2.27948836e-30]], dtype=float32), indices=array([[20, 12, 8, 5, 6]]))




Traffic Sign Key  
 20 : Dangerous curve to the right  
 12 : Priority road  
 8 : Speed limit (120km/h)  
 5 : Speed limit (80km/h)  
 6 : End of speed limit (80km/h)

workzone.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 3.10213411e-10, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00]], dtype=float32), indices=array([[12, 25, 0, 1, 2]]))




Traffic Sign Key  
 12 : Priority road  
 25 : Road work  
 0 : Speed limit (20km/h)  
 1 : Speed limit (30km/h)  
 2 : Speed limit (50km/h)

YCrossing.png




Top five: TopKV2(values=array([[ 1.00000000e+00, 2.00721084e-09, 2.39146881e-27, 1.22845360e-32, 7.74104907e-34]], dtype=float32), indices=array([[12, 26, 20, 25, 11]]))




Traffic Sign Key  
 12 : Priority road  
 26 : Traffic signals  
 20 : Dangerous curve to the right  
 25 : Road work  
 11 : Right-of-way at the next intersection

yeild.png



Top five: `TopKV2(values=array([[ 1., 0., 0., 0., 0.]], dtype=float32), indices=array([[13, 0, 1, 2, 3]]))`



Traffic Sign Key

13 : Yield

0 : Speed limit (20km/h)

1 : Speed limit (30km/h)

2 : Speed limit (50km/h)

3 : Speed limit (60km/h)

yield\_pedestrian.png